

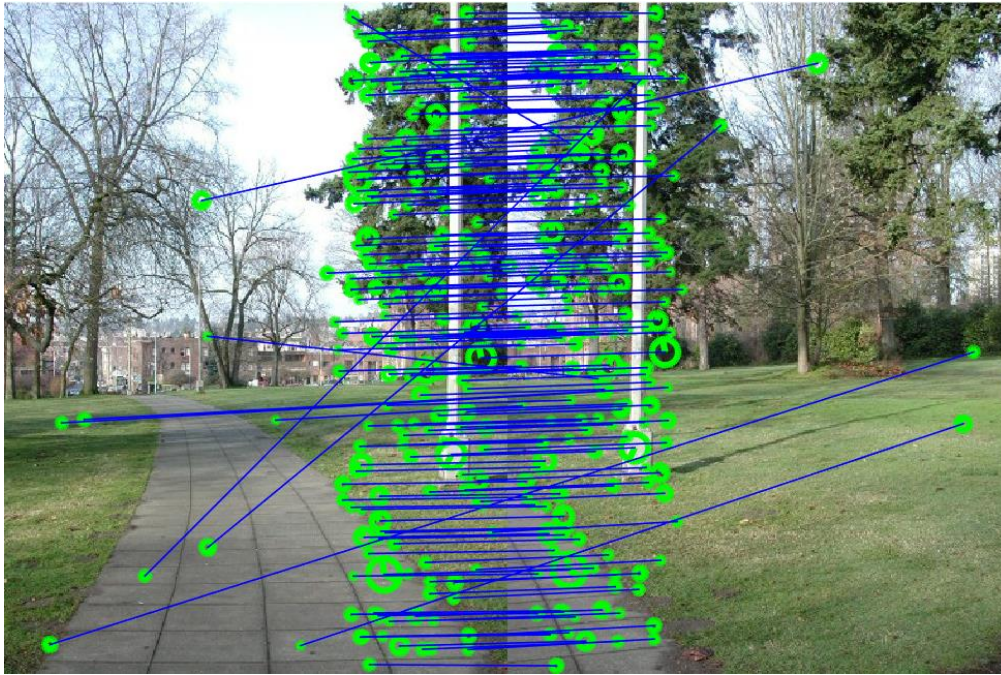
CSE185

Introduction to Computer Vision

Lab 09: Image Stitching

Instructor: Daniel Leung
TA: Mohammadkazem Ebrahimpour
Xueqing Deng

Image Stitching



4 Steps for Image Stitching

1. Feature Extraction (SIFT or Harris)
2. Feature Matching
3. Image Alignment with RANSAC
4. Image Blending/Stitching

4 Steps for Image Stitching

1. **Feature Extraction (SIFT or Harris)**
2. Feature Matching
3. Image Alignment with RANSAC
4. Image Blending/Stitching

SIFT Features Detector

- Use Difference-of-Gaussian to detect multi-scale features

σ doubles for
the next octave

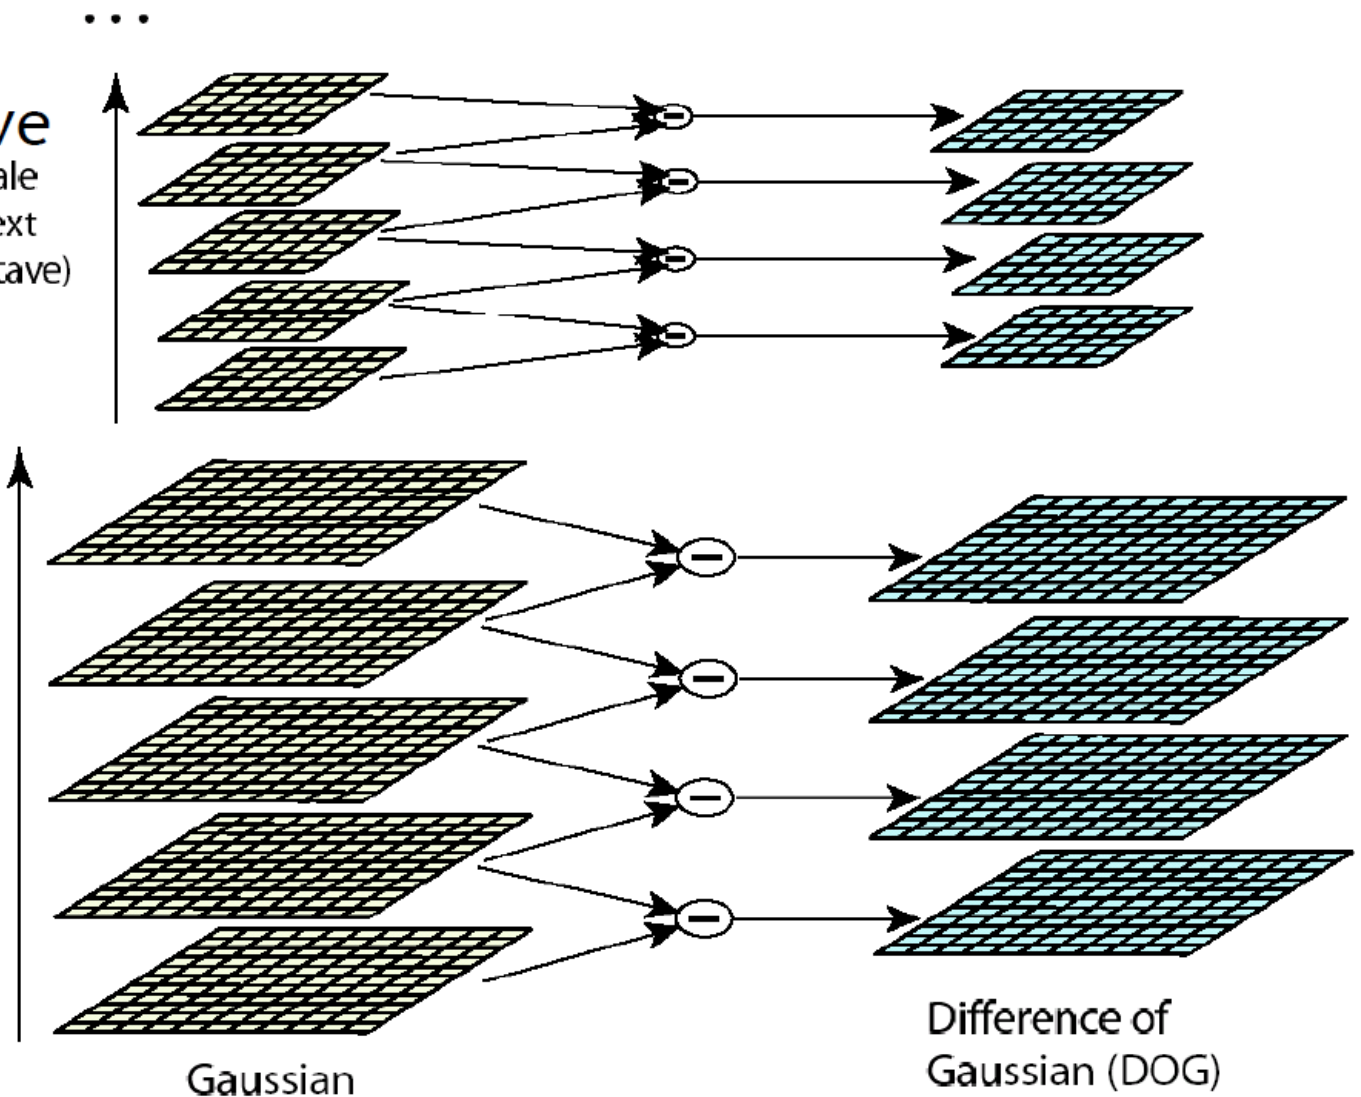
...
Scale
(next
octave)

$$K=2^{(1/s)}$$

Scale
(first
octave)

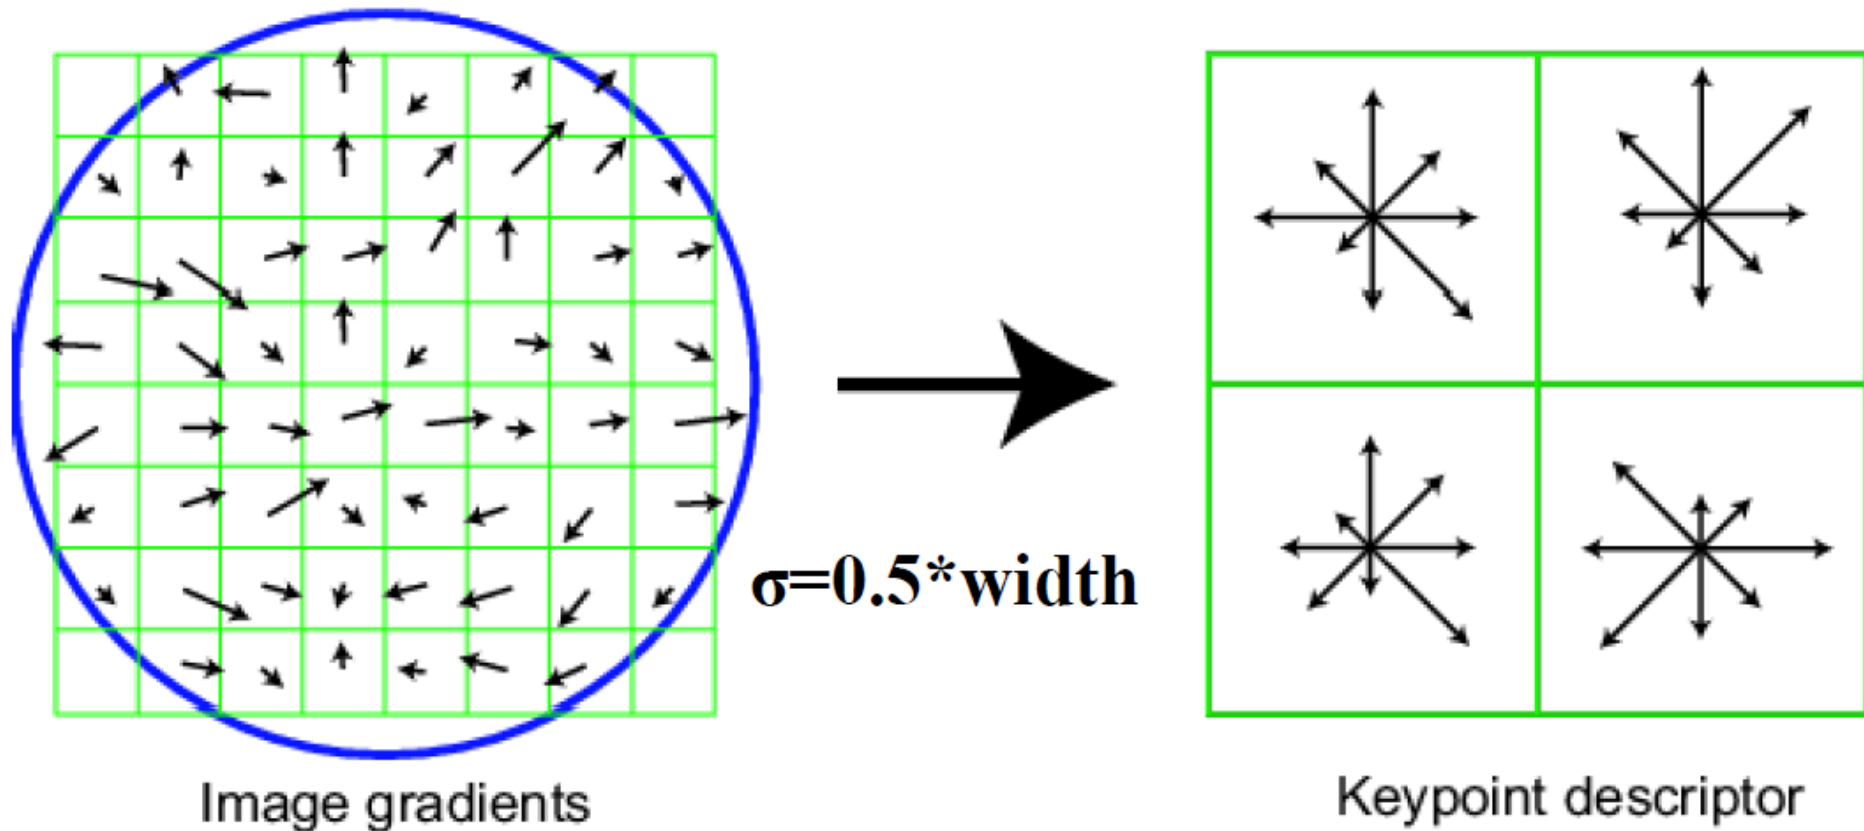
Gaussian

Difference of
Gaussian (DOG)



SIFT Features Descriptor

- Use magnitude and orientation of gradients to describe feature points as 128-D feature vectors



vlfeat

- Use [vlfeat](#) toolbox to extract SIFT features
 - The `vlfeat-0.9.20-bin` folder is already in the archived lab 9 file. After unzipping the lab 9 file, set MATLAB to work within the lab 9 folder.
 - Setup in MATLAB:

```
run('vlfeat-0.9.20-bin/toolbox/vl_setup');
```

- SIFT feature in vlfeat:
<http://www.vlfeat.org/overview/sift.html>

Extract SIFT Features

1. Load image as single format
2. convert to gray scale
3. apply `vl_sift()`

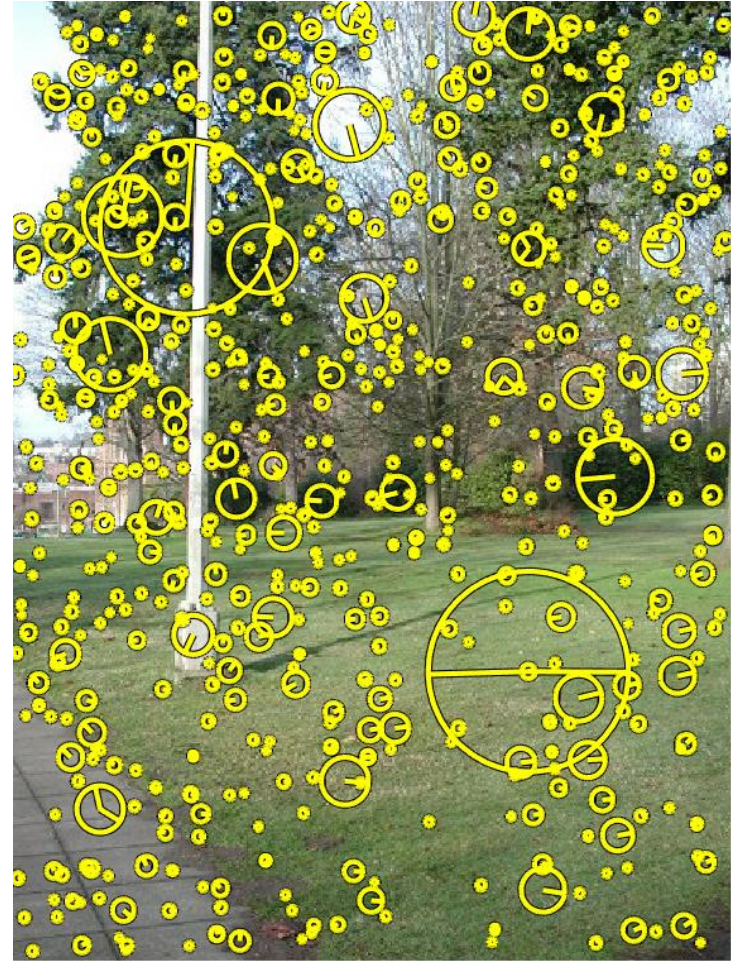
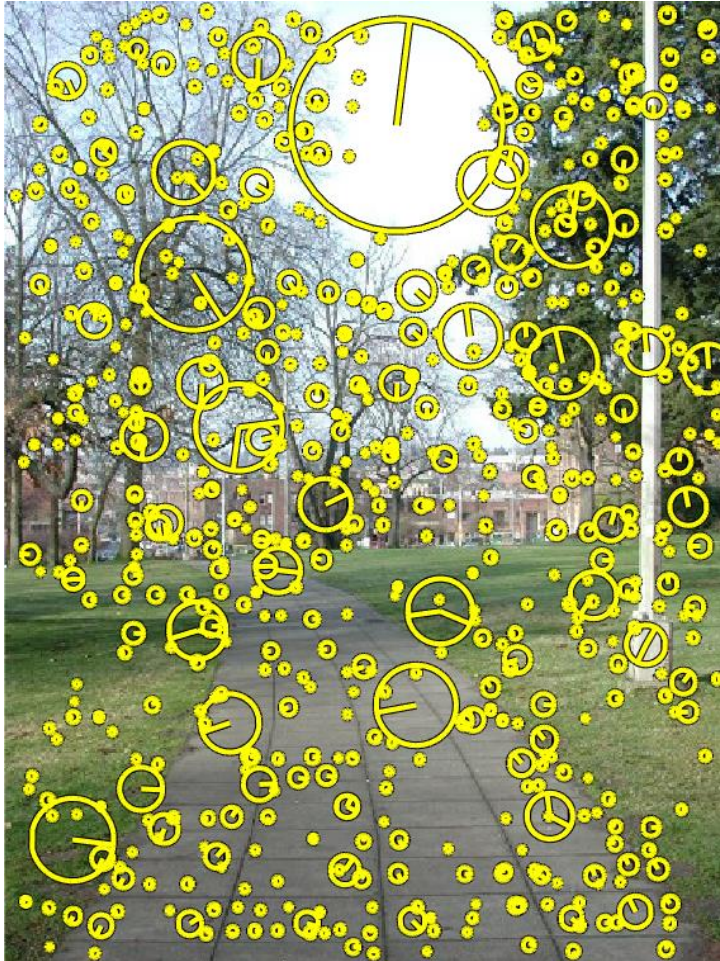
```
img1 = im2single(imread('prtn13.jpg'));  
img2 = im2single(imread('prtn12.jpg'));  
  
%% SIFT feature extraction  
I1 = rgb2gray(img1);  
I2 = rgb2gray(img2);  
  
[f1, d1] = vl_sift(I1);  
[f2, d2] = vl_sift(I2);  
  
d1 = double(d1);  
d2 = double(d2);
```

f is a $4 \times N$ feature frame, each
column = $[x, y, scale, orientation]$

d is a $128 \times N$ descriptor, each
column is an 128-D feature vector

Extract SIFT Features

- Draw feature points with `plot_sift(img, f, d)`

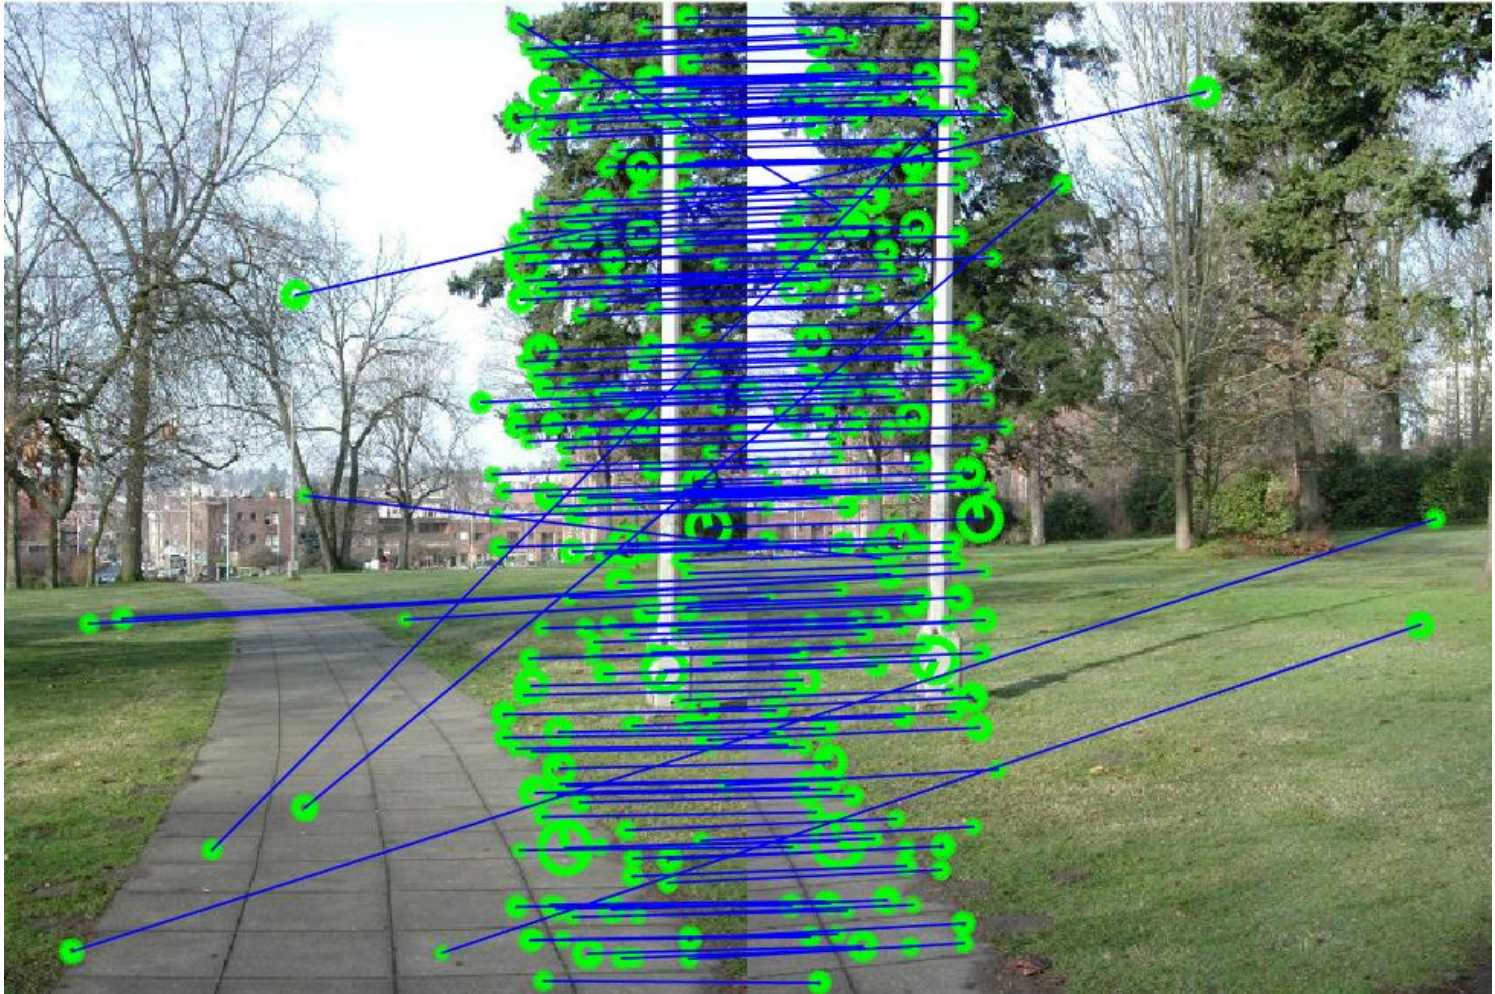


4 Steps for Image Stitching

1. Feature Extraction (SIFT or Harris)
- 2. Feature Matching**
3. Image Alignment with RANSAC
4. Image Blending/Stitching

Feature Matching

- Find matched features between two images



Feature Matching

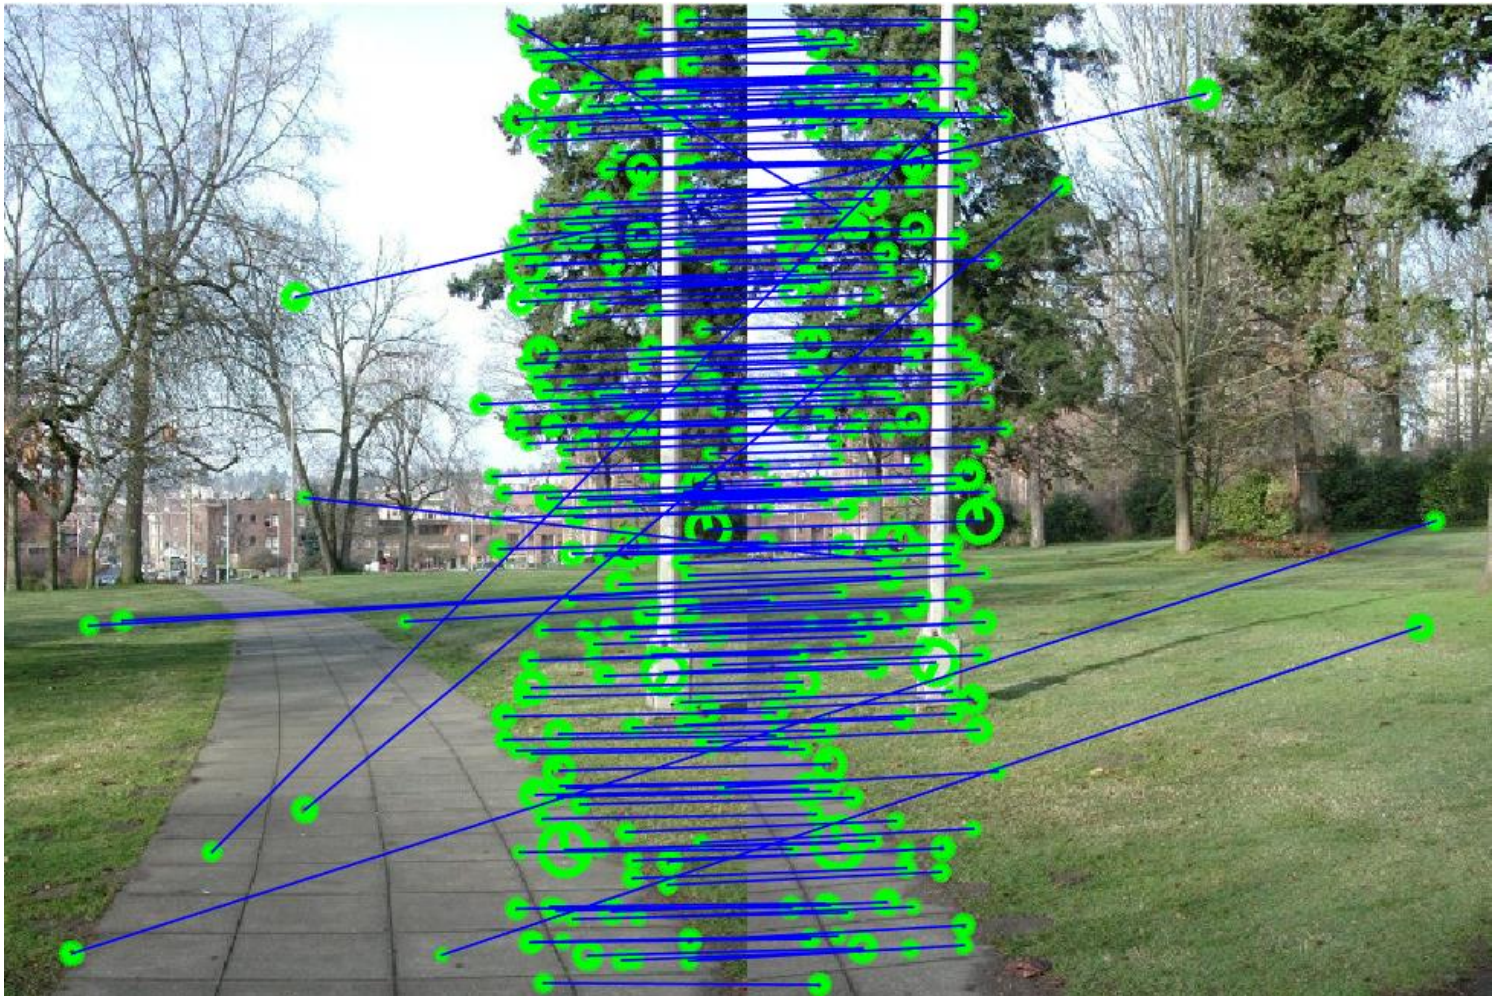
- Use `vl_ubcmatch()`

```
[matches, scores] = vl_ubcmatch(d1, d2) ;
```

- `matches` is a $2 \times N_{match}$ matrix, each column indicates the matched index of `d1` and `d2`
 - `d1(:, matches(1, 1))` and `d2(:, matches(2, 1))` are the first matched pair
 - `d1(:, matches(1, k))` and `d2(:, matches(2, k))` are the k -th matched pair
- Plot matched points with `plot_match(img1, img2, f1, f2, matches)`

Feature Matching

- `plot_match(img1, img2, f1, f2, matches)`



4 Steps for Image Stitching

1. Feature Extraction (SIFT or Harris)
2. Feature Matching
3. **Image Alignment with RANSAC**
4. Image Blending/Stitching

Image Alignment

- Assume simple translation model: $p1 = p2 + \begin{pmatrix} tx \\ ty \end{pmatrix}$



image 1




image 2

Image Alignment

- Assume simple translation model: $p1 = p2 + \begin{pmatrix} tx \\ ty \end{pmatrix}$



image 1

image 2

Image Alignment

- Each matching pair can determine a pair of (tx, ty)




Image Alignment

- Use all matches pair to compute average (tx, ty)
 - outliers cause large error



Image Alignment

- Use RANSAC to select the best (tx, ty)



RANSAC

- Assume total N match pairs

Run k times:

1. randomly choose 1 pair (2 feature points)
2. compute tx_0 and ty_0
3. for other $N-1$ pairs:
 - (a) compute tx_1 and ty_1
 - (b) if (tx_1, ty_1) is close to (tx_0, ty_0) ,
count as an inlier

Return the pair (tx_0, ty_0) with the maximal #inliers

How to determine k ?

- n : the number of selected points (in our case, $n = 2$)
- e : the proportion of outliers
- P : the probability of success rate after k trials

$$P = 1 - (1 - (1 - e)^n)^k$$

$$k = \frac{\log(1 - P)}{\log(1 - (1 - e)^n)}$$

- If $n = 2$, $e = 0.1$, $P = 0.999$, then $k \cong 4$

How to count inliers?

- Compute (tx, ty) from feature frame f

```
p1 = f1(1:2, matches(1, i));  
p2 = f2(1:2, matches(2, i));
```

```
tx = p1(1) - p2(1);  
ty = p1(2) - p2(2);
```

f is a $4 \times N$ feature frame, each
column = $[x, y, scale, orientation]$

$matches$ is $2 \times N_{match}$ matrix,
each column = $[f1_{index}, f2_{index}]$

- A pair (tx_1, ty_1) is inlier of (tx_0, ty_0) if

$$(tx_1 - tx_0)^2 + (ty_1 - ty_0)^2 < \delta$$

RANSAC

- Assume total N match pairs

Run k times:

1. randomly choose 1 pair (2 feature points)
2. compute tx_0 and ty_0
3. $\#inlier = 0$
4. for other $N-1$ pairs:
 - (a) compute tx_1 and ty_1
 - (b) if $(tx_1 - tx_0)^2 + (ty_1 - ty_0)^2 < \delta$:
 $\#inlier = \#inlier + 1$

```
if  $\#inlier > \max\_inlier$ :  
    best_tx = tx_0  
    best_ty = ty_0
```

Return (best_tx, best_ty)

use randperm()

4 Steps for Image Stitching

1. Feature Extraction (SIFT or Harris)
2. Feature Matching
3. Image Alignment with RANSAC
4. **Image Blending/Stitching**

Image Blending/Stitching

- Fix image 1, and shift image 2 by (tx, ty)



image 1

image 2

Image Blending/Stitching

- Simply paste image 2 over image 1



Image Blending/Stitching

- Fix image 1, and shift image 2 by (tx, ty)

```
output = zeros(H + ty, W + tx, 3);
output(1:H, 1:W, :) = img1;

for y2 = 1:size(img2, 1)
    for x2 = 1:size(img2, 2)

        y1 = y2 + ty;
        x1 = x2 + tx;

        if( y1 >= 1 && y1 <= H + ty &&
            x1 >= 1 && x1 <= W + tx )
            output(y1, x1, :) = img2(y2, x2, :);
        end
    end
end
end
```

TODO

- Implement lab09.m
 - use vlfeat for feature extraction (5pt) and feature matching (5pt)
 - implement RANSAC (5pt)
- Adjust e and P to see the difference on k and stitching results
- Try other image pairs (in denny.zip), upload at least 3 results and lab09.m (5pt)